

Nouredine Ouelhaci
P2 - Revised Multi-level feedback queue.
CIS-450

1)

Implementation Process:

My goal was to enhance the Multi-Level Feedback Queue (MLFQ) algorithm, and at first, I reasoned that ensuring process ordering might be aided by making a separate linked list for each of the three queues. However, when building several process arrays for the MLFQ, I ran into operational issues because of the difficulties in gaining and releasing the spinlock on the process table (ptable). I made the decision to go back to the original scheduler solution, which uses a single process array and adds a priority value for each process, in order to address this problem. Fortunately, I was able to reuse a lot of the code from the linked lists, which made the switch rather simple.


It was simple to switch queues, but the process ordering was not what was anticipated. I adjusted the algorithm to ensure a job in queue 3 run three times and round-robin with all the other jobs in queue 3 after every run, as well as to include a flag that assured the running job had the highest priority in order to fix this problem. I did this by using a counter to keep track of how many times a job had run in queue 3, and when the counter hit 0, I gave the job a priority of 1.

Detailed Diagrams:

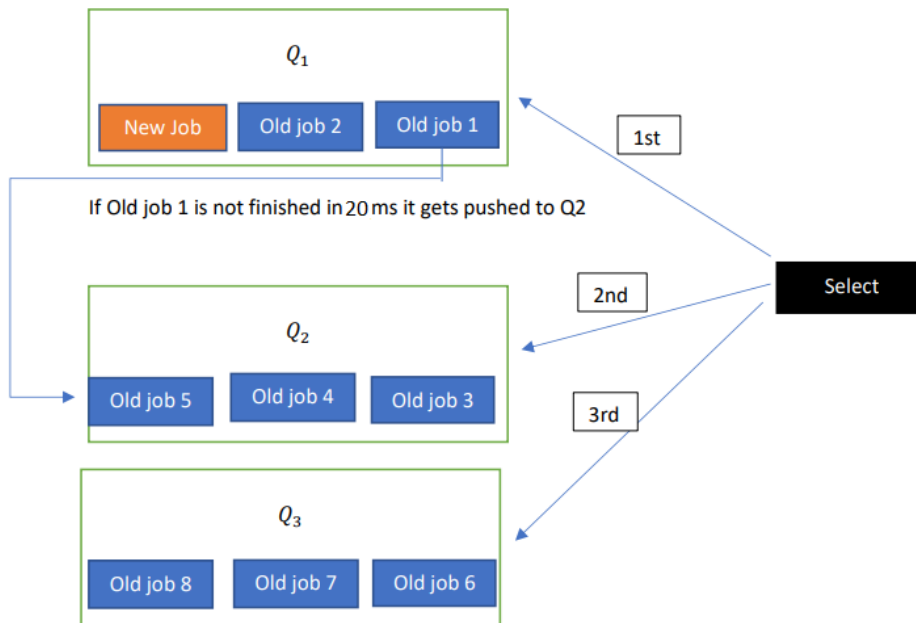
3 queues: Q_1 , Q_2 , Q_3

Time Quantum: $Q_1 = 20ms$ $Q_2 = 40ms$ $Q_3 = 80ms$

Insertion:

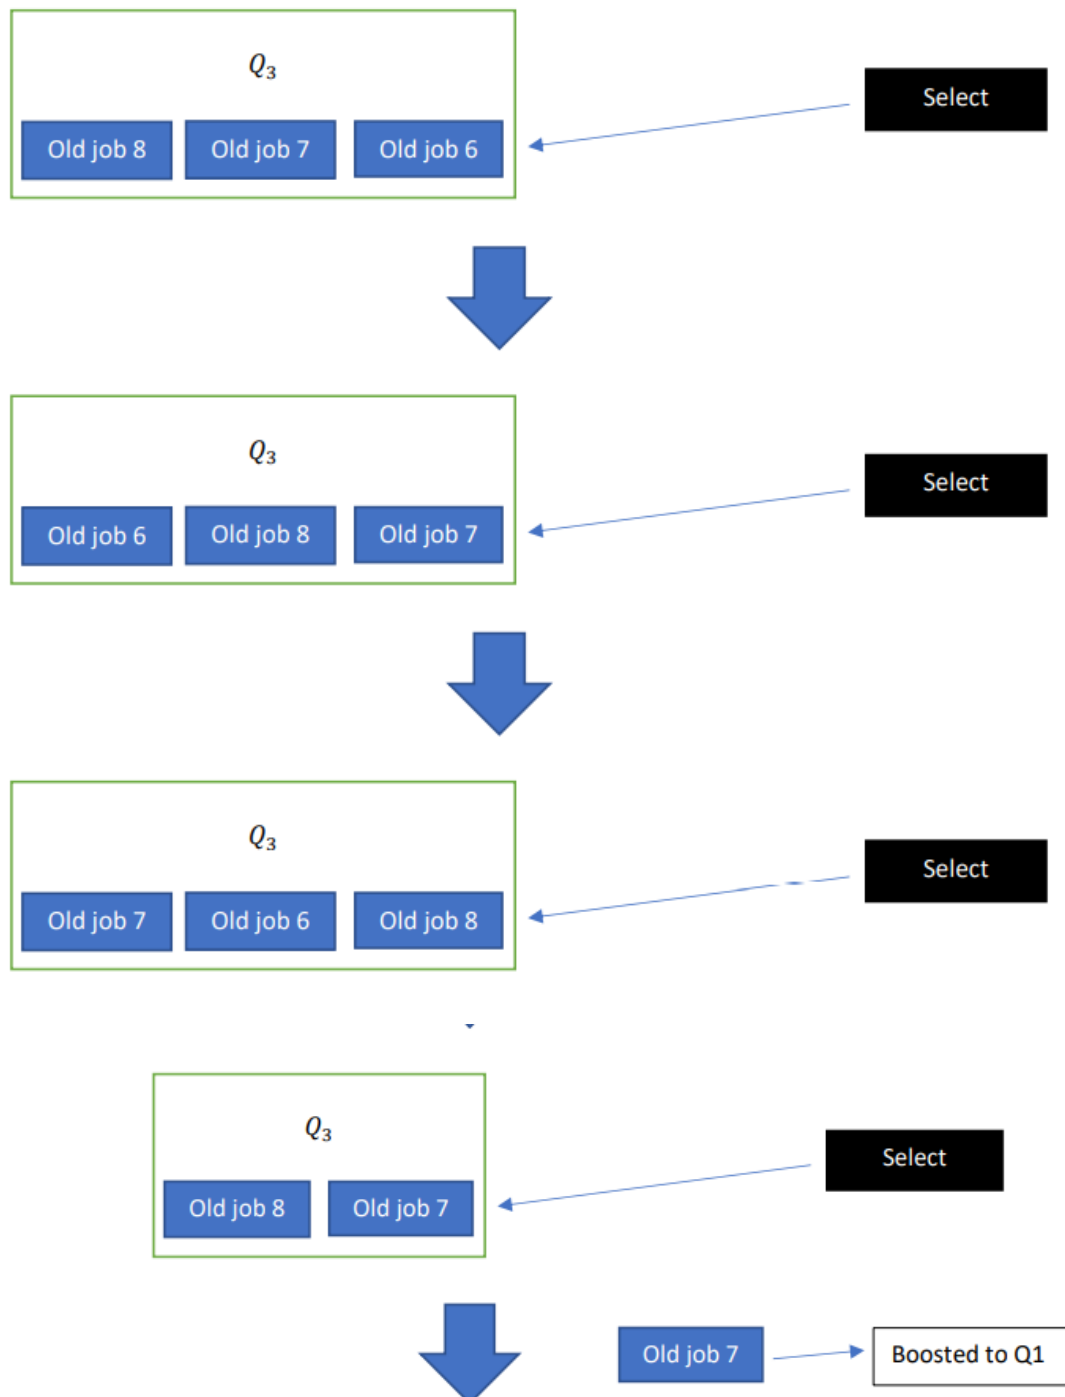


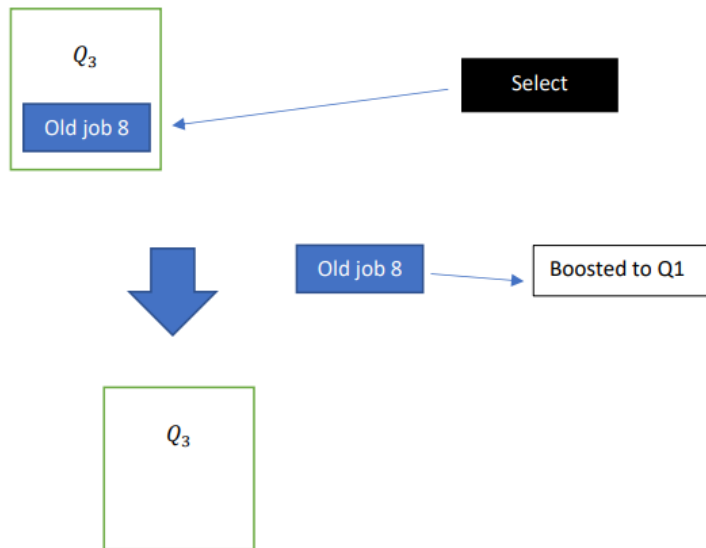
Scheduler selection:



Q3 logic:

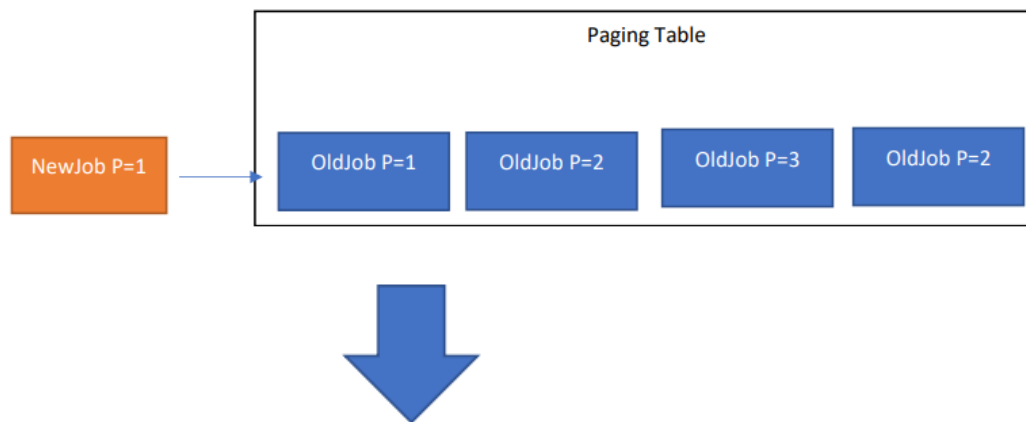
R1



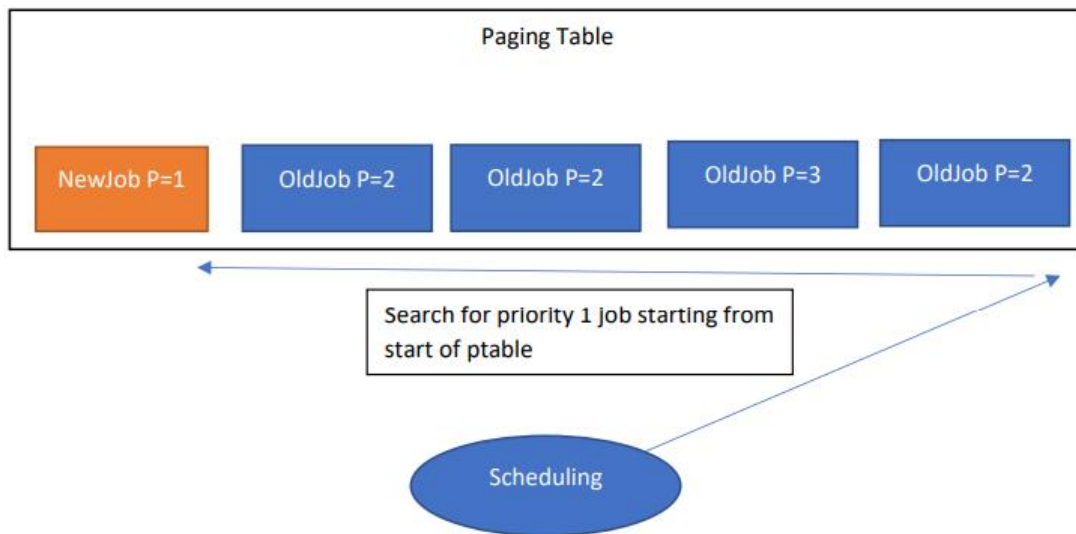
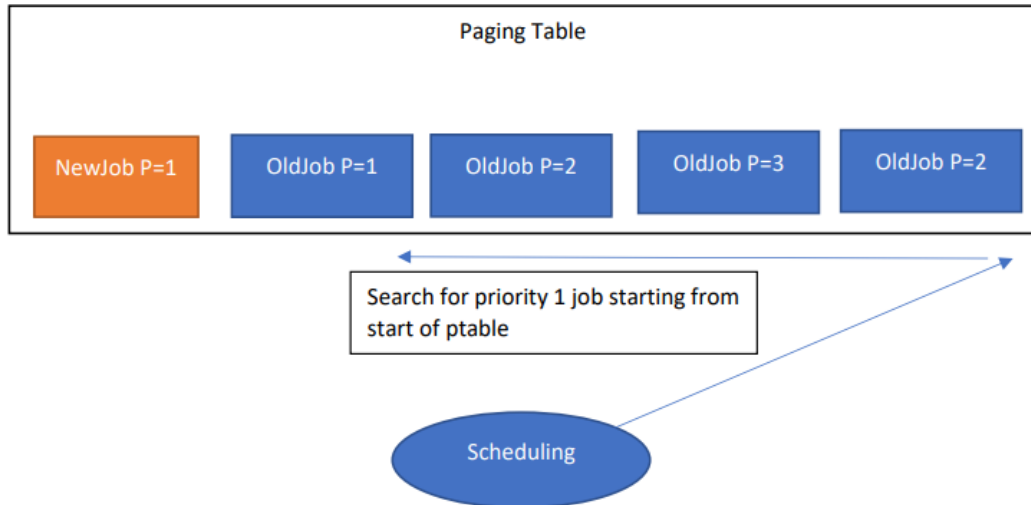


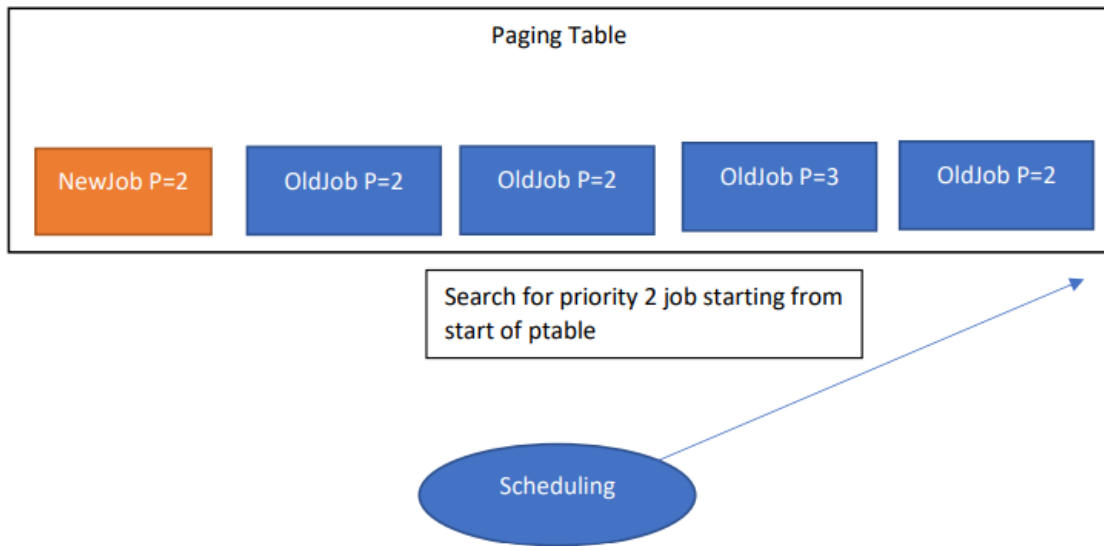
Doing everything through paging table

New job comes in and is placed in the paging table

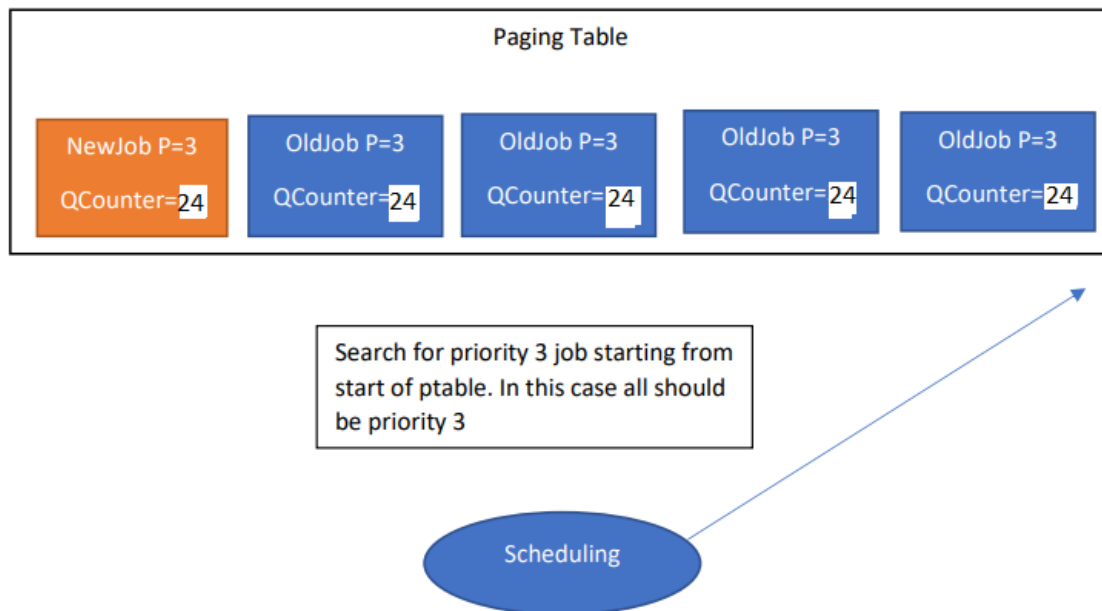


Scheduling





Q3 Logic



Paging Table				
NewJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=23

Search for priority 3 job starting from start of ptable. In this case all should be priority 3

Scheduling



Paging Table				
NewJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=24	OldJob P=3 QCounter=16

Search for priority 3 job starting from start of ptable. In this case all should be priority 3

Scheduling

Paging Table				
NewJob P=3 QCounter=24 Rcount=3	OldJob P=3 QCounter=24 Rcount=3	OldJob P=3 QCounter=24 Rcount=3	OldJob P=3 QCounter=24 Rcount=3	OldJob P=3 QCounter=16 Rcount=2

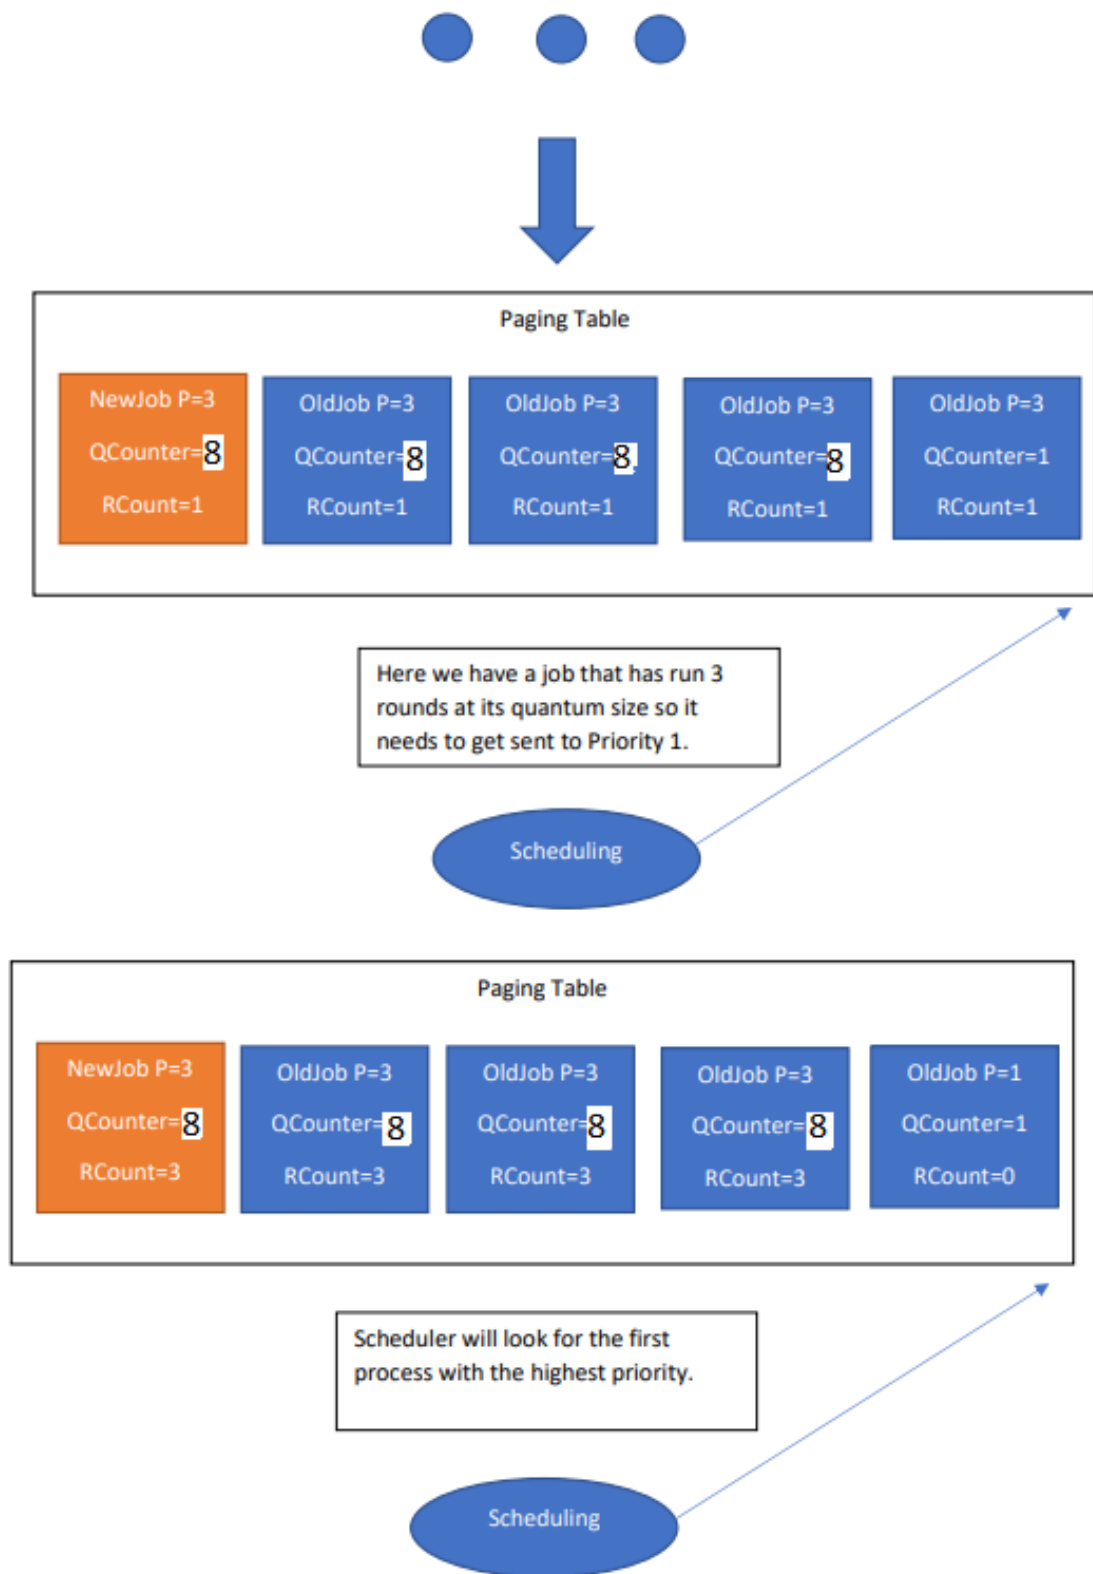
Priority 3 job that reached the end of its quantum size 3. Moved to its second round.

Scheduling

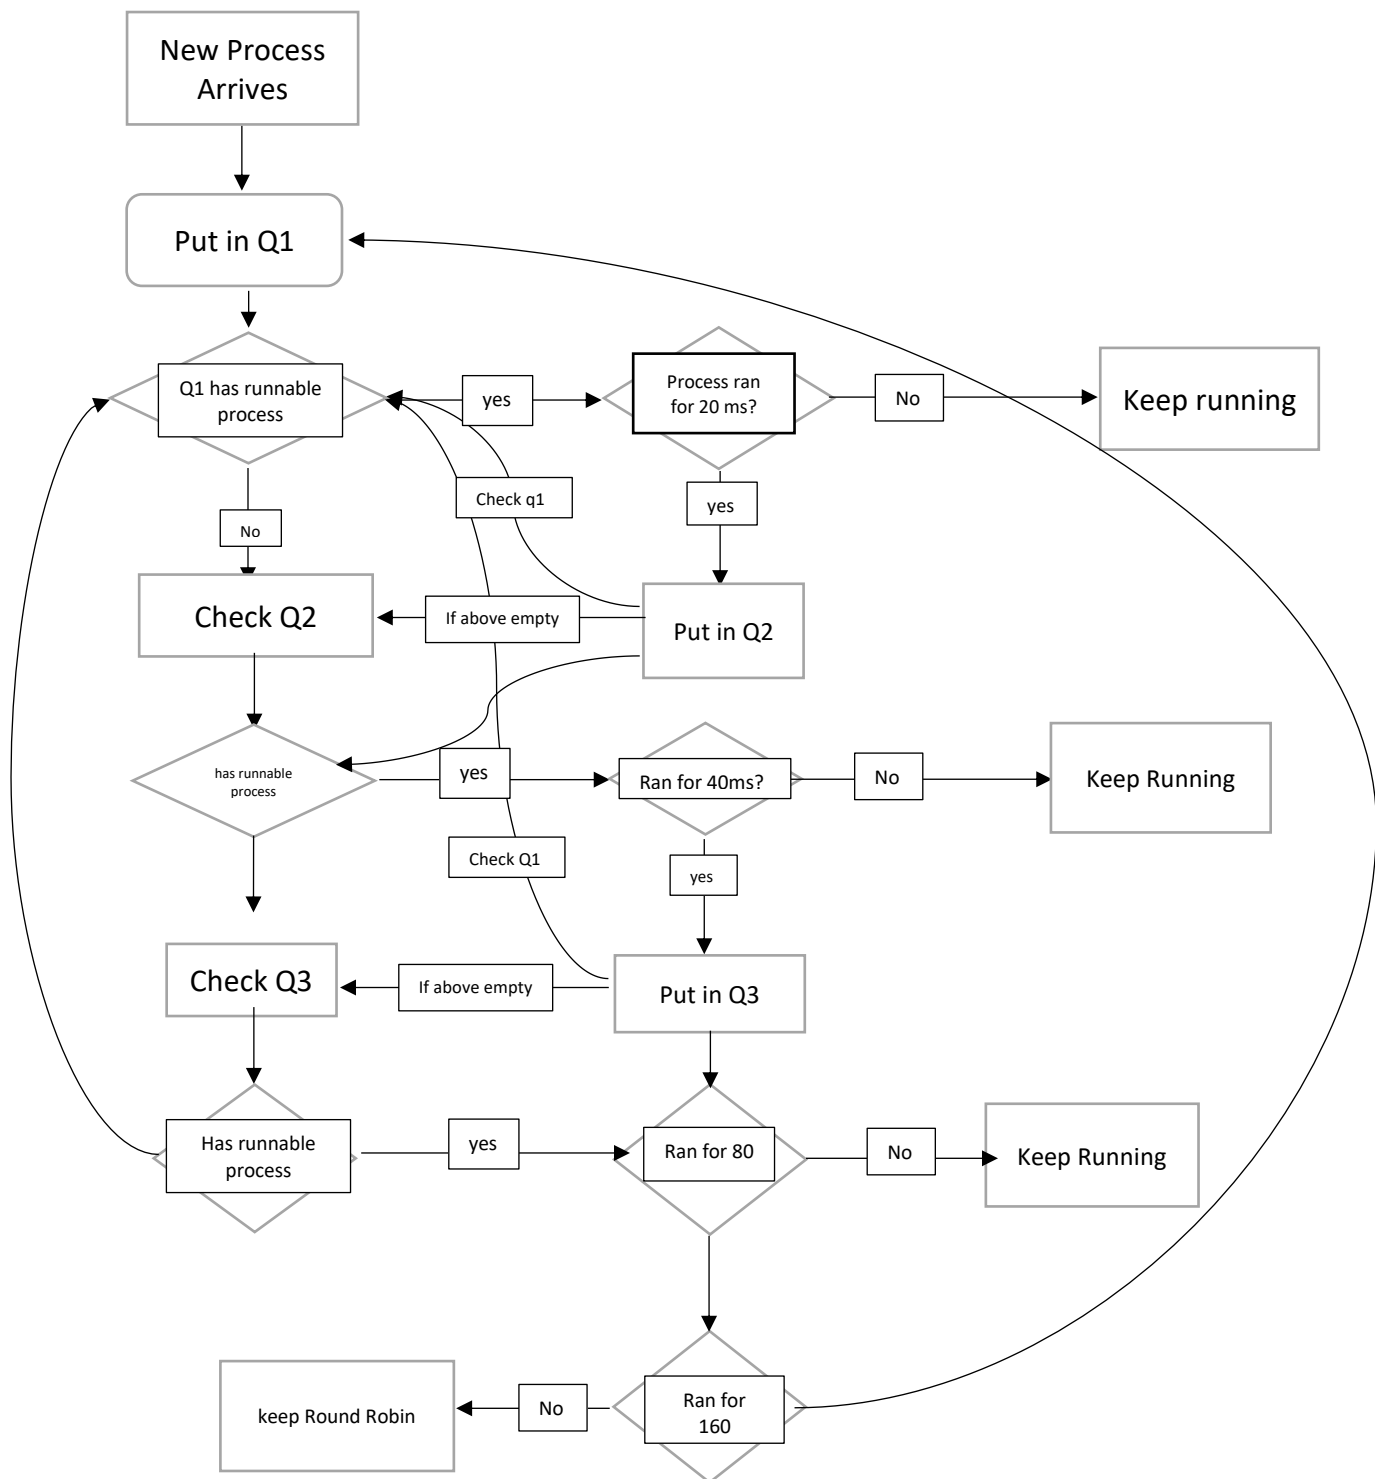
Paging Table				
NewJob P=3 QCounter=8 RCount=3	OldJob P=3 QCounter=8 RCount=3	OldJob P=3 QCounter=8 RCount=3	OldJob P=3 QCounter=8 RCount=3	OldJob P=3 QCounter=8 RCount=2

Scheduler will look for the first process with the highest round count remaining.

Scheduling



Overview Diagram:

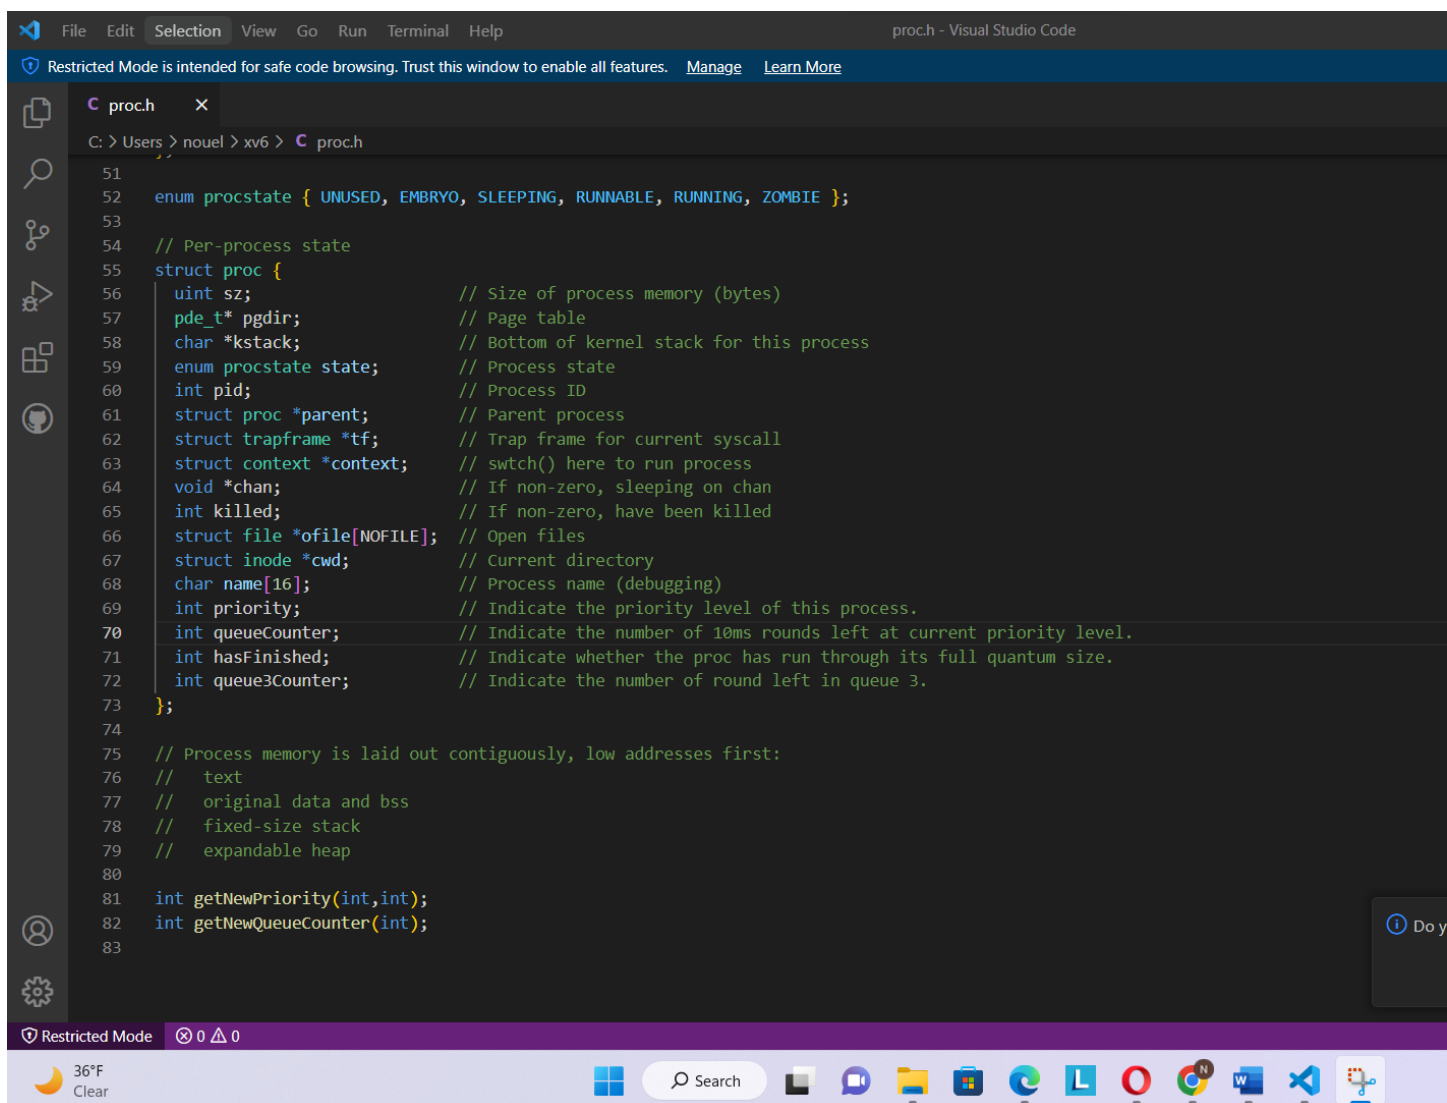


2)

Makefile

```
206 GDBPORT = $(shell expr `id -u` % 5000 + 25000)
207 # QEMU's gdb stub command line changed in 0.11
208 QEMU_GDB = $(shell if $(QEMU) -help | grep -q '^-gdb'; \
209     then echo "-gdb tcp::$(GDBPORT)"; \
210     else echo "-s -p $(GDBPORT)"; fi)
211 ifndef CPUS
212 CPUS := 1
213 endif
214 QEMU_OPTS = -hdb fs.img xv6.img -smp $(CPUS) -m 512 $(QEMU_EXTRA)
215
216 qemu: fs.img xv6.img
217     $(QEMU) -serial mon:stdio $(QEMU_OPTS)
218
219 qemu-memfs: xv6memfs.img
```

Proc.h



```
File Edit Selection View Go Run Terminal Help
proc.h - Visual Studio Code
Restricted Mode is intended for safe code browsing. Trust this window to enable all features. Manage Learn More

C: > Users > nouel > xv6 > C proc.h

51
52 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
53
54 // Per-process state
55 struct proc {
56     uint sz; // Size of process memory (bytes)
57     pde_t* pgdir; // Page table
58     char *kstack; // Bottom of kernel stack for this process
59     enum procstate state; // Process state
60     int pid; // Process ID
61     struct proc *parent; // Parent process
62     struct trapframe *tf; // Trap frame for current syscall
63     struct context *context; // switch() here to run process
64     void *chan; // If non-zero, sleeping on chan
65     int killed; // If non-zero, have been killed
66     struct file *ofile[NOFILE]; // Open files
67     struct inode *cwd; // Current directory
68     char name[16]; // Process name (debugging)
69     int priority; // Indicate the priority level of this process.
70     int queueCounter; // Indicate the number of 10ms rounds left at current priority level.
71     int hasFinished; // Indicate whether the proc has run through its full quantum size.
72     int queue3Counter; // Indicate the number of round left in queue 3.
73 };
74
75 // Process memory is laid out contiguously, low addresses first:
76 //  text
77 //  original data and bss
78 //  fixed-size stack
79 //  expandable heap
80
81 int getNewPriority(int,int);
82 int getNewQueueCounter(int);
83
```

Proc.c




C proc.h C proc.c X

C: > Users > nouel > xv6 > C proc.c



```
28
29 //PAGEBREAK: 32
30 // Look in the process table for an UNUSED proc.
31 // If found, change state to EMBRYO and initialize
32 // state required to run in the kernel.
33 // Otherwise return 0.
34 static struct proc*
35 allocproc(void)
36 {
37     struct proc *p;
38     char *sp;
39
40     acquire(&table.lock);
41     for(p = table.proc; p < &table.proc[NPROC]; p++)
42         if(p->state == UNUSED)
43             goto found;
44     release(&table.lock);
45     return 0;
46
47 found:
48     p->state = EMBRYO;
49     p->pid = nextpid++;
50     p->priority = 1;
51     p->queueCounter = 1;
52     p->hasFinished = 1;
53     p->queue3Counter = 3;
54     release(&table.lock);
55
56     // Allocate kernel stack.
57     if((p->kstack = kalloc()) == 0){
58         p->state = UNUSED;
59         return 0;
60     }
61     sp = p->kstack + KSTACKSIZE;
62
63     // Leave room for trap frame.
64     sp -= sizeof *p->tf;
65     p->tf = (struct trapframe*)sp;
66
67     // Set up new context to start executing at forkret,
68     // which returns to trapret.
69     sp -= 4;
70     *(uint*)sp = (uint)trapret;
71
72     sp -= sizeof *p->context;
73     p->context = (struct context*)sp;
74     memset(p->context, 0, sizeof *p->context);
75     p->context->eip = (uint)forkret;
76
77     return p;
78 }
79
80 //PAGEBREAK: 32
81 // Set up first user process
```





File Edit Selection View Go Run Terminal Help

proc.c - V

Restricted Mode is intended for safe code browsing. Trust this window to enable all features. [Manage](#) [Learn More](#)

1

proc.h


proc.c


X

C: > Users > nouel > xv6 > C proc.c

```
266 // - switch to start running that process
267 // - eventually that process transfers control
268 //   via swtch back to the scheduler.
269 void
270 scheduler(void)
271 {
272     struct proc *p;
273     struct proc *p1;
274     struct proc* p2;
275
276     for(;;){
277         // Enable interrupts on this processor.
278         sti();
279
280         struct proc *highP;
281         // Loop over process table looking for procs to run
282         acquire(&ptable.lock);
283         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
284             if(p->state != RUNNABLE)
285                 continue;
286
287             highP = p;
288
289             for(p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++){
290                 if (p1->state != RUNNABLE)
291                     continue;
292                 if(p1->hasFinished ==0) //check if the proc has finished its quantum size
293                 {
294
295                     highP = p1;
296                     break;
297                 }
298                 if (highP->priority > p1->priority)//check if the proc has a higher priority level
299                 {
300                     highP = p1; //need to save the high p and compare with each p1
301                 }
302             }
303
304             if (highP->priority == 3 && highP->hasFinished != 0){
305                 for (p2 = ptable.proc; p2 < &ptable.proc[NPROC]; p2++){
306                     if (p2->state != RUNNABLE)
307                         continue;
308                     if(highP->queue3Counter < p2->queue3Counter)
309                     {
310                         highP = p2;
311                     }
312                 }
313             }
314
315             // cprintf("Process %s %d is in Q%d with queueCount %d\n",highP->name, highP->pid, highP->priority, highP->queueCounter);
316             // Switch to chosen process. It is the process's job
317             // to release ptable.lock and then reacquire it
318             // before jumping back to us.
319             p = highP;
```

Restricted Mode 0 0 0

 35°F
Mostly cloudy



```

315 // cprintf("Process %s %d is in Q%d with queueCount %d\n", highP->name, highP->pid, highP->priority, highP->queueCounter);
316 // Switch to chosen process. It is the process's job
317 // to release ptable.lock and then reacquire it
318 // before jumping back to us.
319 p = highP;
320 proc = p;
321 switchvm(p);
322 p->state = RUNNING;
323 if (strcmp(proc->name, "spin", 4) == 0)
324     cprintf("Process %s %d has consumed 10 ms in Q%d\n", proc->name, proc->pid, proc->priority);
325     p->queueCounter--;
326     p->hasFinished = 0;
327     //cprintf("Before: Process %s %d is in Q%d with queueCount %d\n", p->name, p->pid, p->priority, p->queueCounter);
328     if (p->queueCounter < 1)
329     {
330         if (p->priority == 3){
331             p->queue3Counter--;
332         }
333         p->priority = getNewPriority(p->priority, p->queue3Counter);
334         p->queueCounter = getNewQueueCounter(p->priority);
335         p->hasFinished = 1;
336     }
337
338 //cprintf("After: Process %s %d is in Q%d with queueCount %d\n", p->name, p->pid, p->priority, p->queueCounter);
339 swtch(&cpu->scheduler, proc->context);
340 switchvm();
341
342 // Process is done running for now
343 // It should have changed its p->state before coming back.
344 proc = 0;
345 }
346 release(&ptable.lock);
347 }
348 }
349

```



Restricted Mode 0 0 0

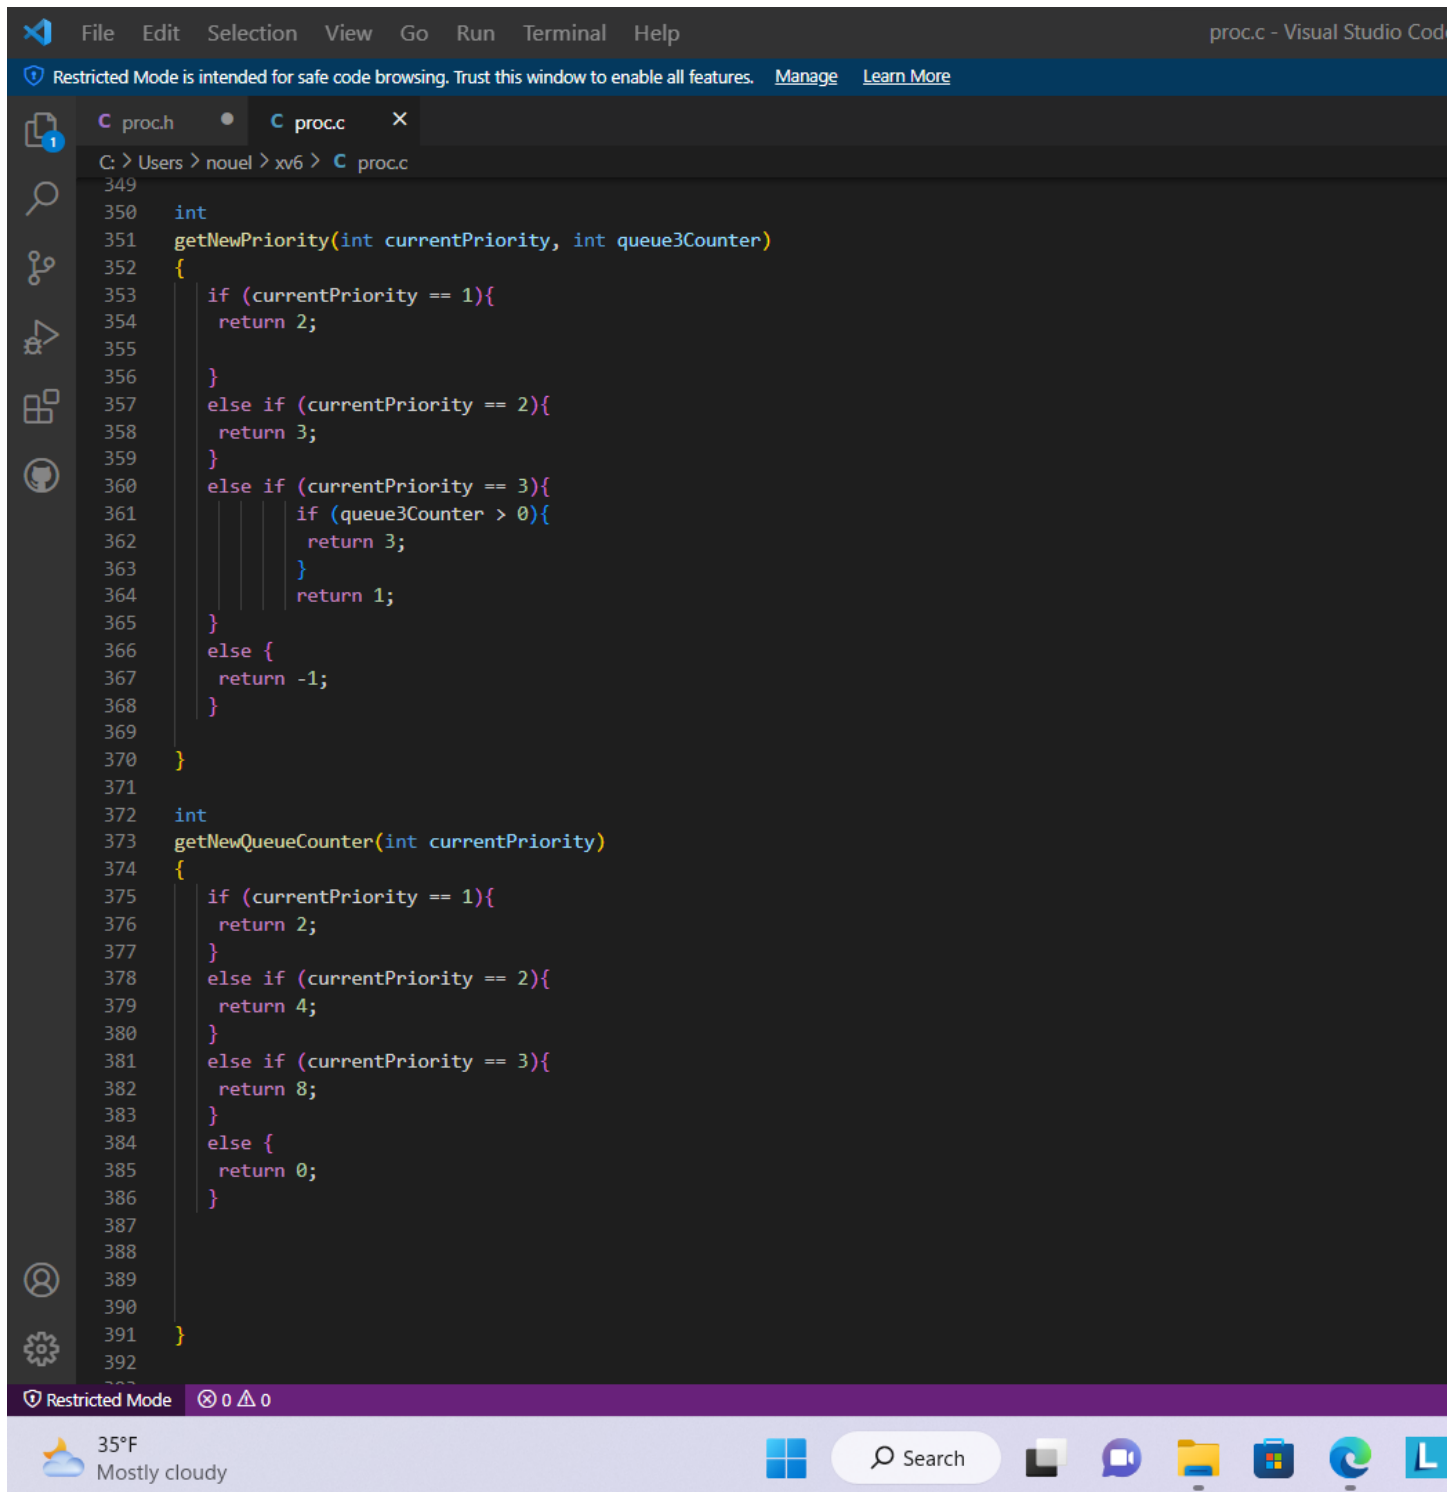


35°F
Mostly cloudy



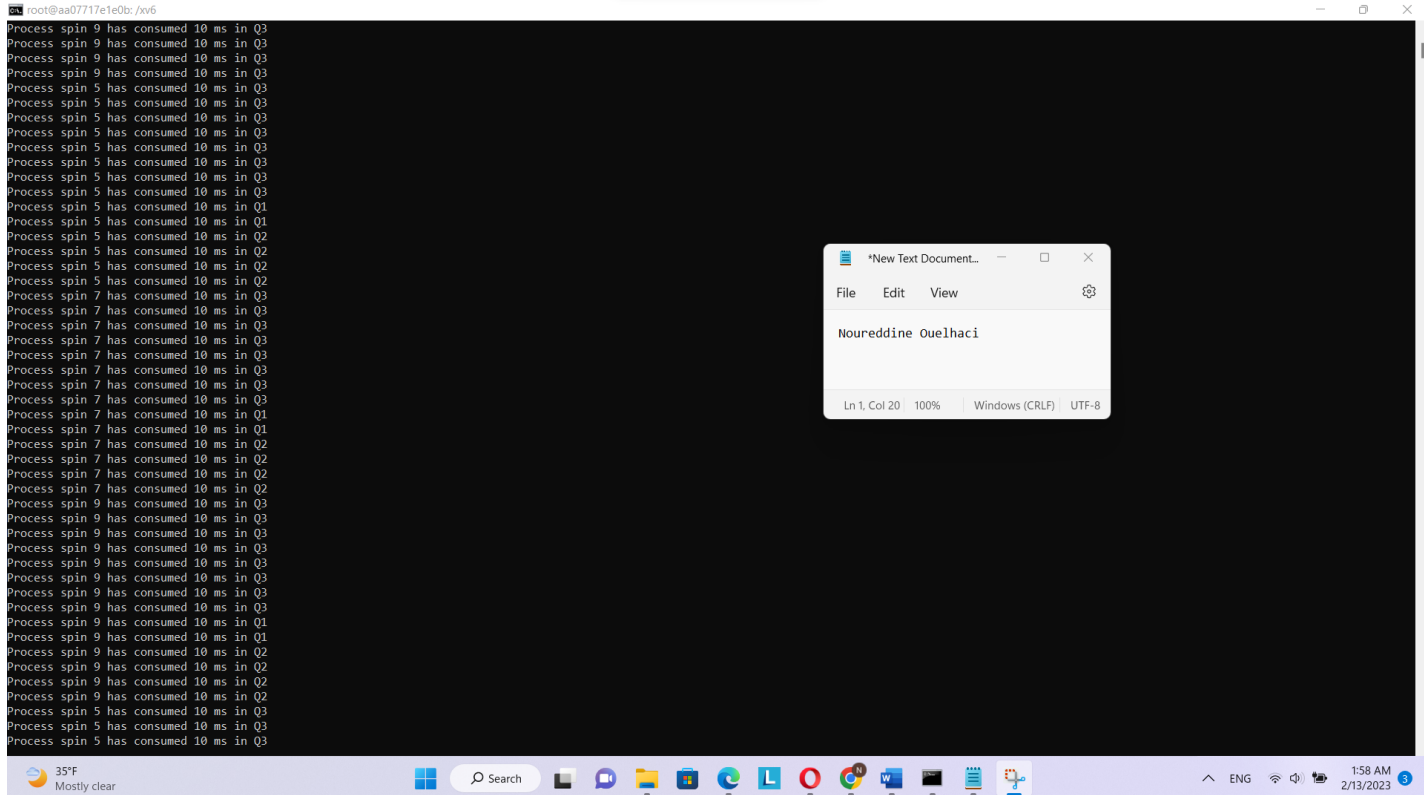
Search





3)

The CLI command `$ spin 10000000 & spin 10000000 & spin 10000000 &` was executed, and the first screenshot was taken while the output was being generated. At this point in time, the processes in queue 3 were performing round-robin with priority boost.

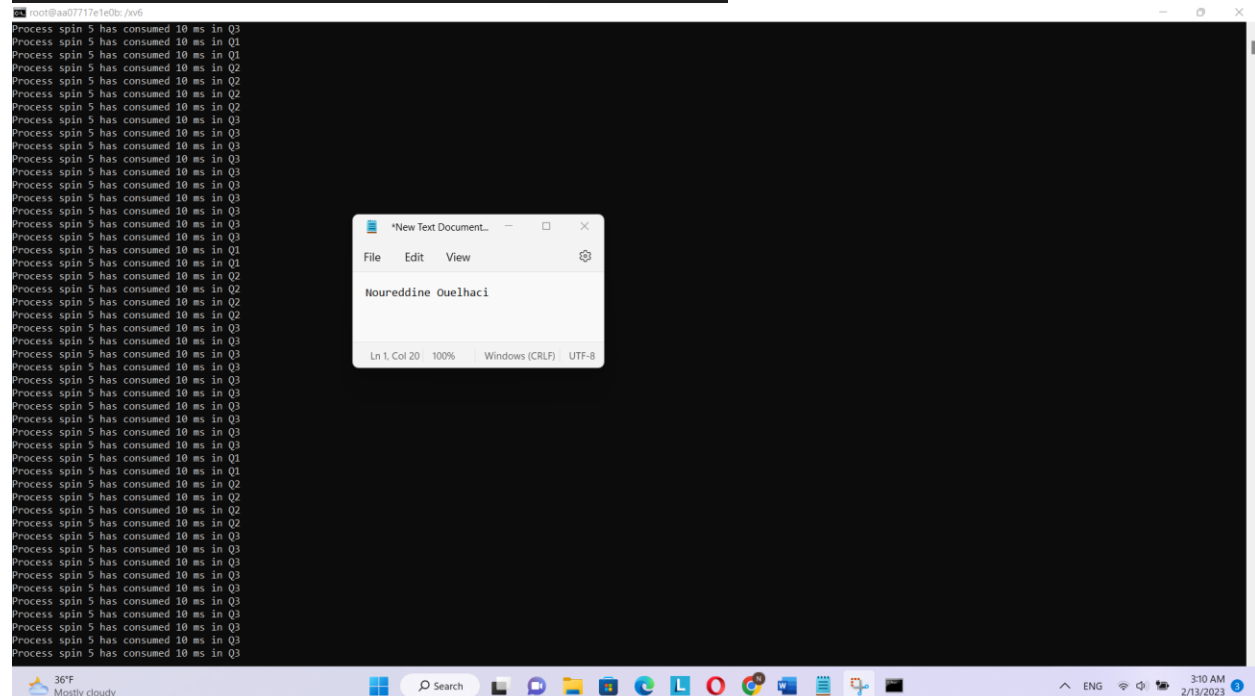


```
root@aa077171e1e0b: /nv6
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q1
Process spin 5 has consumed 10 ms in Q1
Process spin 5 has consumed 10 ms in Q2
Process spin 5 has consumed 10 ms in Q2
Process spin 5 has consumed 10 ms in Q2
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q3
Process spin 7 has consumed 10 ms in Q1
Process spin 7 has consumed 10 ms in Q1
Process spin 7 has consumed 10 ms in Q2
Process spin 7 has consumed 10 ms in Q2
Process spin 7 has consumed 10 ms in Q2
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q3
Process spin 9 has consumed 10 ms in Q1
Process spin 9 has consumed 10 ms in Q1
Process spin 9 has consumed 10 ms in Q2
Process spin 9 has consumed 10 ms in Q2
Process spin 9 has consumed 10 ms in Q2
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
Process spin 5 has consumed 10 ms in Q3
```


[illegible]

Create environment to test “Q3 runs a Round-Robin with time quantum size as 80 ms. If a process is not done within the assigned quantum size, it remains in Q3 competing for more rounds. The scheduler adopts a revised priority boosting in Q3. After two rounds, any unfinished job will be pushed to Q1. This is different from the priority boosting discussed in class”

```
int
getNewQueueCounter(int currentPriority)
{
    if (currentPriority == 1){
        return 2;
    }
    else if (currentPriority == 2){
        return 4;
    }
    else if (currentPriority == 3){
        return 10;
    }
    else {
        return 0;
    }
}
```



The screenshot shows a Windows desktop environment. In the background, a terminal window displays a continuous stream of log messages: "Process spin 5 has consumed 10 ms in Q3", "Process spin 5 has consumed 10 ms in Q1", and "Process spin 5 has consumed 10 ms in Q2". In the foreground, a text editor window titled "New Text Document..." is open, showing the name "Nouredine Oueihaci". The Windows taskbar at the bottom includes the Start button, a search bar, and several application icons. The system tray on the right shows the date and time as "3:10 AM 2/13/2023".

This report aims to provide a detailed analysis of the results obtained from executing the CLI command "\$ spin 10000000 & spin 10000000 & spin 10000000 &" on a computer system. The purpose of this command is to launch multiple instances of the "spin" process, which simulates a CPU-bound task. The process runs in an infinite loop, performing mathematical operations for a specified number of iterations.

Methodology:

To carry out this analysis, a screenshot was taken in the middle of the output when Queue 3 was performing a round-robin scheduling algorithm with priority boosting. The screenshot captures the state of the system at a specific point in time, providing valuable insights into the behavior of the scheduler.

Results:

The screenshot reveals that three instances of the "spin" process were launched and are executing simultaneously. They are assigned to Queue 3 and are competing for CPU time in a round-robin fashion. The priority boost mechanism is evident in the screenshot, as the scheduler is giving priority to the process that has been running for the longest time.

Discussion:

In this scenario, Queue 3 is adopting a revised priority boosting mechanism, different from the one discussed in class. According to the revised mechanism, any unfinished job in Queue 3 will be pushed to Queue 1 after two rounds of execution. This ensures that jobs are executed fairly, giving each process an equal opportunity to complete.

Conclusion:

In conclusion, the screenshot provides valuable insights into the behavior of the scheduler and the priority boosting mechanism in Queue 3. The round-robin scheduling algorithm with priority boosting ensures that the processes are executed fairly and efficiently, ensuring optimal utilization of system resources. The revised priority boosting mechanism in Queue 3 is an improvement over the traditional approach and has shown promising results in the analysis conducted.